


Legend:

- C1. cudaMalloc w/ caching (vanilla Pytorch)
- A2. Application-level model weights offload
- U0. UVM w/o caching (Pytorch default)
- U1. Pytorch-UVM



OPT model size (2nd line: Oversubscription ratio)