


C1. cudaMalloc w/ caching (vanilla Pytorch) U0. UVM w/o caching (Pytorch default)
A. Application-level KV Cache Offload U1. Pytorch-UVM



of concurrent prompts
(2nd line: Oversubscription ratio)