

- C1. cudaMalloc w/ caching (vanilla Pytorch)
- A. Application-level KV Cache Offload
- U0. UVM w/o caching (Pytorch default)
- U1. Pytorch-UVM
- U3. Pytorch-UVM + no migration

