


Total Time (sec)

C1. cudaMalloc w/ caching (vanilla Pytorch) U0. UVM w/o caching (Pytorch default)
A. Application-level KV Cache Offload



of concurrent prompts
(2nd line: Oversubscription ratio)