

■ C1. cudaMalloc w/ caching (vanilla Pytorch) ■ U0. UVM w/o caching (Pytorch default)
■ A. Application-level KV Cache Offload

