

■ C1. cudaMalloc w/ caching (vanilla Pytorch)    ■ U0. UVM w/o caching (Pytorch default)  
■ A. Application-level KV Cache Offload

