


C1. cudaMalloc w/ caching (vanilla Pytorch) U0. UVM w/o caching (Pytorch default)
A2. Application-level model weights offload U1. Pytorch-UVM



OPT model size (2nd line: Oversubscription ratio)