


■ C1. cudaMalloc w/ caching (vanilla Pytorch)
 ■ U0. UVM w/o caching (Pytorch default)
 ■ A2. Application-level model weights offload
 ■ U1. Pytorch-UVM



Qwen model (2nd line: Oversubscription ratio)